

DETERMINISTIC CORE - 0 LLM tokens, reproducible, unit-testable

Project data
tasks · people
deps · comments

Knowledge graph
typed, in-memory

NetworkX

Risk metrics
6 graph algorithms

critical path, artic. pts

Witness subgraph
k=1, ≤60 nodes

O(anomalies)

only the witness subgraph crosses this boundary

risk scores bypass the LLM entirely

NARRATION LAYER - bounded prompt; never sees the raw dataset

LLM narration
explains ONE finding

cites node ids only

Grounding check
drop_ungrounded()

citation in subgraph?

Verified output
risk + evidence

persisted as AgentRun

fabricated citations stripped (200/200)