

TeamSync AI - IEEE Paper Reference

Contents

1. One-page summary
2. Problem statement
3. Proposed solution
4. Novelty — the four claims
5. Architecture
6. Features
7. Datasets
8. Evaluation — all metrics
9. Limitations & threats to validity
10. Figures
11. Related work positioning
12. IEEE paper structure mapping
13. Reproducibility
14. What is still missing

Explainable, Graph-Grounded Multi-Agent Project Intelligence

This is the single consolidated reference for writing the paper. Every number is measured, reproducible, and independently re-audited. Nothing is estimated. 4. Novelty · 6. Features · 8. Evaluation · 10. Figures · 12. IEEE structure · 14. Gaps

1. One-page summary

What it is	A Kanban project-management platform with an AI layer that detects six kinds of team coordination risk, explains each with verifiable evidence, and recommends actions.
Core idea	Don't ask an LLM to find risk. Compute risk deterministically from a knowledge graph, then let the LLM only narrate one finding from a bounded slice of that graph.
Main claim	Risk detection can be made both cheap (487× fewer tokens) and trustworthy (mechanically verified citations) without losing accuracy.
Headline numbers	487× token reduction · held-out F1 0.712 beats a tuned ML model's 0.593 · ablation F1 1.000 vs 0.947 vs naive prompting · κ 0.95 LLM-vs-deterministic agreement · 200/200 fabricated citations stripped
Real dataset	TAWOS — 458,232 real Jira issues, 39 open-source projects (MSR 2022, Apache 2.0)
Stack	Next.js + FastAPI ×2 + PostgreSQL + NetworkX + Groq/LLM
Suggested venue	IEEE Access (good fit now) · regional IEEE conferences (comfortably above bar)

2. Problem statement

Teams using Jira / Linear / Asana hit four compounding failures:

1. **Risk is detected late.** Tools show current state, not risk trajectory. An overloaded engineer or a slipping critical-path task is visible only after the damage.
2. **Warnings are opaque.** When a tool flags "at risk," it rarely shows *which evidence* produced that verdict — so a manager can neither act on it nor challenge it.
3. **Signals are siloed.** No single view reveals a cross-cutting risk such as *"the only person who understands this component is also unresponsive and on the critical path."*
4. **Naive AI fixes create new problems.** Feeding raw project data to an LLM is expensive (cost scales with project size), non-deterministic (same input, different verdicts), and produces justifications nobody can verify.

Research question. *Can a system detect coordination risk early, explain it with verifiable evidence, and stay cheap enough to run continuously — without trading away accuracy?*

3. Proposed solution

Six steps. The inversion in steps 2–4 is the whole idea.

1. Compile project state into a **typed knowledge graph** (people, tasks, milestones, dependencies, comments) — from data the product already stores.
2. Compute all six risks as **deterministic graph algorithms** — zero LLM tokens, zero randomness.
3. For each detected anomaly, extract a small **witness subgraph** (only the relevant nodes).
4. Give the LLM **only that subgraph**, and ask it to *narrate* — never to decide.
5. **Mechanically verify** every citation against the subgraph; discard anything that doesn't resolve to a real node.
6. Every agent has a **rule-based fallback** with an identical schema, so the system works with no LLM at all.

The LLM never sees or judges the raw dataset. It only explains an already-computed, already-verified fact.

Problem	How this solves it
Late detection	Graph algorithms are cheap enough to run on every change, not as a periodic report
Opaque warnings	Every finding carries evidence that is mechanically checkable
Siloed signals	One graph unifies task, people, dependency and comment data
Expensive/unreliable AI	Prompt is bounded; the numeric answer never depends on the LLM

4. Novelty — the four claims

Each is independently testable, and each was tested.

Novelty 1 — Graph-Grounded Agent Prompting

Project state → typed knowledge graph → risk computed as **graph invariants** (critical path, articulation points, weighted degree, community structure), not inferred by a language model. The LLM receives one finding plus its minimal witness subgraph and must cite node ids from it. This inverts the standard "LLM finds the problem" pattern and is what makes claims 2–4 possible.

Novelty 2 — Prompt cost is $O(\text{anomalies})$, not $O(\text{project size})$

Because the LLM only ever sees a bounded subgraph, prompt size stops tracking project size. **Measured: 266 → 390 tokens (1.47×) while the project grows 154× (13 → 2,003 tasks).** Against naive full-dataset prompting: **189,834 → 390 tokens = 487× reduction.**

Novelty 3 — Hallucinated evidence is mechanically detectable

Most XAI systems ask an LLM to cite sources and hope. Here every citation is checked against a bounded set of real node ids; anything unresolvable is stripped. **Measured: 200/200 fabricated references removed.** "Told to cite evidence" becomes "every surviving citation is verifiably real."

Novelty 4 — A zero-parameter graph rule beats a tuned ML model on unseen data

Three attempts to improve accuracy with more sophisticated methods **all failed or backfired** on held-out projects. The untuned deterministic rule scored **F1 0.712** against the tuned logistic model's **0.593** on 22 projects neither had seen. Reported as a finding in its own right: added model complexity was actively harmful here.

One-sentence novelty statement (for the abstract):

Risk detection can be made both cheap (graph-computed, $O(\text{anomalies})$ prompt cost) and trustworthy (mechanically grounded evidence) without sacrificing accuracy to a more complex learned model.

5. Architecture

5.1 System topology

Three services; the AI service is independently deployable and degrades to a rule-based path without the LLM.




Figure 1

Service	Tech	Port	Role
Frontend	Next.js 14, React 18, Tailwind, TS	3000	UI, Kanban, dashboards, WebSocket
Core API	FastAPI, SQLAlchemy 2.0 async, Pydantic v2	8000	Auth, CRUD, multi-tenancy, persistence
AI Service	FastAPI, NetworkX, Groq client	8001	Graph build, risk metrics, agent pipeline
Database	PostgreSQL 16 (SQLite in tests)	5432	System of record
Cache/Queue	Redis 7 + Celery	6379	Caching, background tasks

5.2 The analysis pipeline — where the contribution lives




Figure 2

Key property: steps A–D and the risk scores involve **no LLM call at all**. Only narration touches the model, and only on a bounded slice.

5.3 Knowledge graph schema

Nodes: Person, Task, Milestone, Project, Component

Edge	Derived from
Person -[ASSIGNED_TO {points}]-> Task	task assignee + story points
Person -[REPORTED]-> Task	task reporter
Task -[BLOCKS]-> Task	task dependency
Task -[SUBTASK_OF]-> Task	parent task
Task -[PART_OF]-> Milestone	milestone link
Person -[COMMENTED_ON]-> Task	task comments
Person -[MEMBER_OF]-> Project	project membership
Task -[TOUCHES]-> Component	milestone-as-component proxy
Person -[KNOWS {depth}]-> Component	ASSIGNED_TO o TOUCHES

All edges come from existing tables — **no new data collection required.**

5.4 The six risks as graph algorithms

Risk type	Graph computation	Threshold
Dependency (critical path)	Longest path through BLOCKS DAG	$\text{CRITICAL_PATH_SHARE} \geq 0.25$
Dependency concentration	Betweenness centrality	—
Knowledge / SPOF	Articulation points, Person↔Component	$\text{SPOF_MIN_COMPONENT_TASKS} = 2$
Workload	Weighted degree over ASSIGNED_TO	$\text{WORKLOAD_SKEW} \geq 1.5$
Coordination	Community structure / isolated members	—
Silent member	COMMENTED_ON + ASSIGNED_TO degree	$\text{SILENT_RATIO} \geq 0.30$
Delay	Open tasks past due date	$\text{OVERDUE_RATIO} \geq 0.10$

Free correctness win: the pre-existing API allowed creating a 3-hop circular dependency ($A \rightarrow B \rightarrow C \rightarrow A$) because it only checked direct self-loops. `nx.simple_cycles` catches real cycles once the graph exists.

5.5 Witness subgraph — the efficiency mechanism

Parameter	Value	Why
DEFAULT_HOPS	1	Only immediate neighbourhood of an anomaly
MAX_NEIGHBORS_PER_SEED	8	Critical — without it, a high-degree node re-introduces size scaling
MAX_WITNESS_NODES	60	Hard ceiling; seeds always survive truncation

Serialization is **columnar**, not JSON: field names appear once per block header (`tasks(id,status,pts,due,title):`) instead of on every row.

5.6 Agent pipeline

Coordinator → 3 specialists (parallel) → Risk aggregator → Recommendation generator

Agent	Reads	Produces	Status
Planning	Milestones, capacity, graph metrics	Sprint readiness, capacity gaps	Active
Progress	Tasks, velocity, graph metrics	Velocity trend, stalled work	Active
Workload	Assignments, graph metrics	Overload, SPOFs	Active
Risk	The 6 graph-computed scores	Aggregated scores + evidence	Active

Agent	Reads	Produces	Status
Recommendation	Risk + specialist outputs	Prioritized actions	Active
Meeting Intel	Transcripts	Decisions, action items	Deferred — no data source
Comm Intel	Chat/email events	Response delays, gaps	Deferred — no data source
Review Trio	A proposal	Consensus verdict	Deferred

Shared output schema for **every** agent (LLM and fallback alike):

```
AgentOutput:
  summary, risk_level[low|medium|high|critical], confidence[0-1],
  signals[{name, value, weight}],
  evidence[{source, reference_id, excerpt, relevance}],
  recommendations[{type, title, description, reasoning, priority, confidence}],
  next_action, metadata # metadata.fallback flags the deterministic path
```

6. Features

Product: organization multi-tenancy · JWT auth with refresh rotation · projects with milestones and health/risk scores · Kanban board (drag-drop, dependencies, subtasks, comments) · status workflow Backlog→Planned→In Progress→Blocked→Review→Done · WebSocket real-time updates.

AI layer: 5 active agents covering 6 risk types · deterministic zero-token risk scoring · optional LLM narration · rule-based fallback for every agent (works with no API key) · mechanically enforced evidence grounding · full audit trail (every run persists specialist outputs, risk scores, and the witness subgraph used) · real multi-hop cycle detection.

Deliberately deferred (documented scope, not hidden gaps): Meeting Intelligence, Communication Intelligence, Review Trio, Organizational Memory/RAG, Team Intelligence Index, fine-grained RBAC, persistent graph storage (Neo4j).

7. Datasets

7.1 TAWOS — the real-world dataset

Property	Value
Full name	A Versatile Dataset of Agile Open Source Software Projects
Citation	Tawosi, Al-Subaihin, Moussa & Sarro, MSR 2022, doi:10.1145/3524842.3528029
License	Apache 2.0 (citation required)
Content	Real Jira issue-tracker data, 12 public Jira repositories
Size	4.3 GB MySQL dump

Verified row counts (0 parse errors across the full dump, reconciled against published totals):

Table	Rows
project	39
user	206,162
sprint	4,594
issue	458,232
issue_link	246,587
comment	1,518,327

How it is used — sprint-level delay prediction. TAWOS has no `due_date` field, so the label is derived: a sprint is *delayed* when $\geq 30\%$ of its issues were unresolved at sprint end.

Point-in-time reconstruction. Every issue in an archival dataset is long since resolved, so current status cannot be used — that leaks the future. Each task's state is rebuilt *as of sprint end*, and metrics are evaluated with `now = sprint_end + 1 day`.

Eligible sprints (CLOSED, has end date, ≥ 15 issues): **987 sprints, 31 projects, 58.8% positive.**

Split discipline — by *project*, not sprint, so a team's Jira conventions cannot leak across the boundary. Held-out set evaluated **exactly once**.

Split	Projects	Sprints
Tuning	9	717
Held-out	22	270 (48.9% positive)

7.2 Synthetic scenarios

180 generated scenarios (90 positive / 90 negative, 30 per risk type) with a **deterministically injected** anomaly, so ground truth is known by construction. Used for implementation-correctness checks, boundary sweeps, and the ablation.

7.3 Datasets considered but not used

Dataset	For	Status
AMI Meeting Corpus	Meeting Intelligence	Deferred — agent not built
Enron Email Corpus	Communication Intelligence	Deferred — agent not built

8. Evaluation — all metrics

8.1 Engineering baseline

Item	Value
Backend tests	29/29 pass
AI service tests	43/43 pass (pytest -m mvp)
Frontend tests	27/27 pass (vitest)
Code size	~21K LOC
Quality gates	ruff, mypy, eslint, prettier, tsc

8.2 Synthetic detection correctness

180 scenarios, anomaly injected deterministically.

Metric	Value
TP / FP / FN / TN	90 / 0 / 0 / 90
Precision / Recall / F1 / Accuracy	1.00 / 1.00 / 1.00 / 1.00
Stability	Identical across 4 further seeds (1,200 extra trials)

State this honestly in the paper. For deterministic threshold code on unambiguous inputs, 100% is the *expected* result of correctness — equivalent to "all unit tests pass," **not** evidence of generalization.

Boundary sweep (the genuinely falsifiable check — does each detector flip exactly at its documented threshold?):

Signal	Documented	Observed flip
overdue_ratio	0.10	exactly 0.10
silent_ratio	0.30	0.20 → 0.30 (exact)
workload_skew	1.50	bracketed 1.429 → 1.522

8.3 Efficiency — the headline result

Full two-series sweep, same measurement method for both:

Project size (tasks)	Graph-grounded	Naive full-dataset	Ratio
13	266	1,050	3.9×
53	386	4,695	12.2×
203	389	18,901	48.6×
1,003	389	94,724	243.5×
2,003	390	189,834	486.8×

Derived metric	Value
Token growth vs size growth	1.47× tokens for 154× project size
Reduction at 2,003 tasks	99.8% (487×)

Latency (deterministic path: build + metrics + select; 203 tasks, n=50):

p50	p95	mean	max
10.96 ms	13.42 ms	11.87 ms	50.58 ms

Varies with machine load between runs (p50 observed 4.8–11.0 ms across runs). The claim is "single- to low-double-digit milliseconds" — which is what justifies running analysis **synchronously**, without a queue.

Grounding enforcement: 200/200 fabricated references stripped (100%).

8.4 Real-world results — TAWOS

All 987 sprints (default, untuned):

Metric	GraphMetrics	Story-point baseline
Precision	0.630	0.545
Recall	1.000	0.386
F1	0.773	0.452
Accuracy	0.655	0.450
AUC-ROC	0.581	—
Brier	0.346	—

Held-out — 270 sprints, 22 unseen projects (THE headline number):

Model	Precision	Recall	F1	Accuracy	AUC
Graph rule (untuned, 0 parameters)	0.552	1.000	0.712	0.604	—
Tuned composite logistic model	0.620	0.568	0.593	0.619	0.658

Confusion matrices — rule: TP 132, FP 107, FN 0, TN 31 · composite: TP 75, FP 46, FN 57, TN 92.

Interpretation: the system catches **every** delayed sprint (recall 1.00) while over-flagging roughly half the on-time ones. It is a **high-sensitivity early warning**, not a precise classifier — the right operating point when a false alarm is cheap to dismiss but a missed one is not.

8.5 Negative results — three failed improvement attempts

Reported prominently. They establish the simple rule isn't leaving accuracy on the table.

#	Attempt	Outcome
1	Threshold sweep (0.05–0.50)	Provably inert — byte-identical output at every value. With one shared sprint-end due date, <code>overdue_ratio</code> is always exactly 0.0 or 1.0, so no threshold in (0,1) can change any classification.
2	Individualized due dates from cycle time (2.96 days/story-point)	Un-collapsed the signal (median ratio 0.76) but lowered AUC 0.581 → 0.531. Story points are too noisy a duration proxy.
3	Composite logistic regression, 4 features, 3-fold group CV by project	Won on tuning (CV F1 0.816, AUC 0.887), lost on held-out (F1 0.593). Dominant weight fell on <code>idle_member_ratio</code> , which is mechanically coupled to the label rather than independently predictive; calibration didn't transfer. AUC collapse 0.887→0.658 is the signature of overfitting to project-specific structure.

Conclusion: the zero-parameter, fully explainable rule outperforms the tuned learned model on unseen projects. Added complexity was not merely unhelpful — it was **actively harmful**.

8.6 Ablation — does graph-grounding cost accuracy?

Both architectures scored on **identical** scenarios against **identical** ground truth (n=36, 18 pos / 18 neg, 6 per risk type; **0 failed calls**).

Architecture	Precision	Recall	F1	Accuracy
Graph-grounded (0 tokens, deterministic)	1.000	1.000	1.000	1.000
Naive full-dataset prompting	0.900	1.000	0.947	0.944

Per risk type — **identical on five of six**; the entire gap is coordination:

Risk type	Graph F1	Naive F1
dependency	1.00	1.00
knowledge	1.00	1.00
workload	1.00	1.00
delay	1.00	1.00
silent_member	1.00	1.00
coordination	1.00	0.75 (P 0.60)

The design deliberately favours the naive path, so a graph win can't be dismissed as a strawman: small projects (10–20 tasks, its best case), full context in one call, good-faith prompt defining all six risk types, and lenient response parsing.

What this does NOT show (three honest qualifications):

- The naive baseline is **not bad** — F1 0.947 is respectable. It reports 2.75 risk types per scenario vs the graph's 2.42, with only **12.1%** of its claims uncorroborated. It over-reports only modestly.
- Scoring is **per-target-risk-type**, so extra risks reported on other dimensions aren't counted against it. Stricter scoring would likely widen the gap but isn't claimed.
- Nothing is shown about naive accuracy **at scale** — untested, and largely untestable on any reasonable budget, which is itself part of the argument.

Token note: mean prompt here was 1,109 (graph) vs 1,804 (naive) — only 1.6×, **not** the headline result. These are deliberately tiny projects. The 487× gap appears at 2,003 tasks. Do not conflate the two.

8.7 Live-LLM metrics — narration layer

Three runs; the third is reported. Model: openai/gpt-oss-120b (Groq deprecated the Llama models originally targeted).

Run	Scenarios	Calls	max_tokens	Pacing	Schema validity	κ (n)
v1	18	54	600	4s	27.8%	0.35 (15)
v2	18	54	2000	18s	61.1%	0.84 (33)
v3 (reported)	30 (all 6 types)	90	2000	20s + 429 backoff	41.1%	0.95 (37)

Cohen's κ = 0.9533 (n=37) — "almost perfect" agreement between the LLM's narrated `risk_level` and the deterministic fallback's verdict on identical input. Direct evidence that **the LLM narrates rather than re-decides**.

Risk type	n	κ
dependency	14	1.00
knowledge	12	0.86
workload	8	1.00
coordination / delay / silent_member	1 each	1.00

Schema validity needs two numbers, not one. Of 24 diagnosed failures, **22 were HTTP 429** rate-limits from the free-tier API key — an infrastructure constraint, not a model property.

Definition	Value
Over all attempted calls (incl. rate-limit blocks)	41.1% (37/90)
Over calls that actually reached the model	~92% (37/40)

Quote the first for "how often does an analysis complete end-to-end on this infrastructure"; the second for "how often does the model produce a valid response." Conflating them misleads.

8.8 Metric glossary

Metric	What it measures
Precision / Recall / F1	Classification quality of the risk detector
Accuracy	Overall correctness (less informative under 59/41 class imbalance)
AUC-ROC	Threshold-independent rank quality of the risk score
Brier score	Calibration of the emitted confidence
Point-biserial correlation	Feature screening before building the composite
Cohen's κ	LLM-vs-deterministic verdict agreement
Token count	Evidence for Novelty 2
Latency p50/p95	Justifies synchronous execution
Grounding rate	Evidence for Novelty 3

8.9 Independent verification

Every number was re-audited from scratch, not re-quoted:

- Confusion-matrix arithmetic recomputed by hand — exact match
- Zero project overlap between splits confirmed; $717 + 270 = 987$ partitions cleanly

- Composite model's normalization stats confirmed frozen from tuning-only fitting
- 5 random held-out sprints' labels cross-checked against **raw SQL**, bypassing the ORM — exact match
- `_auc()` verified against brute-force $O(n^2)$ pairwise comparison + known cases — exact match
- `_brier()` verified against known edge cases — exact match
- Both headline scripts re-run fresh — **byte-identical** output
- SQLite row counts re-queried — match TAWOS published totals exactly

No discrepancy found in any check.

8.10 Engineering rigor — bugs found and fixed

Useful for a "lessons learned" or methodology-credibility paragraph.

Bug	Impact
Witness subgraph expanded all neighbours	Size still scaled with project (4.31× growth); fixed with per-seed cap → 1.007×
Grounding treated empty list as "no filter"	Silently allowed everything through; fixed to <code>is not None</code>
Coordinator passed wrong input type	Every specialist failed on every run, swallowed by <code>gather(return_exceptions=True)</code>
Risk aggregation keyed by agent name, looked up by risk type	Risk output was constant regardless of input
Credentials bound as query params	Passwords in access logs; fixed with <code>OAuth2PasswordRequestForm</code>
Due date == eval time in TAWOS harness	<code>due < now</code> never true → recall 0.0 on first run
Dependency link names guessed	Matched only ~2K of ~18.6K real links
Circular-reasoning sweep caught	A proposed sweep reproduced the label definition exactly (F1=1.00) — caught and discarded

9. Limitations & threats to validity

Must appear in the paper. Do not bury these.

#	Threat
1	The delay label is a proxy — TAWOS has no due dates; "≥30% unresolved at sprint end" is defensible but not canonical.
2	Only 1 of 6 risk types validated on real data. SPOF, cycles and coordination are ~0% non-zero in sprint reconstruction (issue links span a project's whole history). The other five are validated synthetically only.
3	<code>blocked_ratio</code> never fires — TAWOS's status vocabulary has no blocked state.
4	AUC is weak by construction at sprint level, since all tasks in a reconstructed sprint share one deadline.
5	Live-LLM sample is thin and imbalanced — $\kappa=0.95$ rests on $n=37$, with 3 of 6 risk types at $n=1$. Single provider, single model, one rate-limited key.
6	No human evaluation of explanation quality. Grounding proves citations are real, not that explanations are useful.
7	No confidence intervals — held-out F1 0.712 is a single point estimate from one split.
8	RBAC not enforced at fine granularity; the guarantee is authenticated + org-scoped.
9	Two designed agents unimplemented (Meeting, Comm Intel) — no data source exists yet.
10	No longitudinal study — no evidence a flagged risk predicts a real-world outcome over weeks.

10. Figures

Four publication-ready SVGs in `ai-service/eval/figures/`, all **516pt = 7.17in wide** (exact IEEE two-column text width), 7–9pt labels, generated from committed result JSONs by `python -m eval.make_figures` — so a figure cannot drift from its evaluation.

File	Content	Section
<code>fig1_architecture.svg</code>	Deterministic core vs narration layer, with the single crossing boundary	§5
<code>fig2_token_scaling.svg</code>	Log-log token scaling, 487× gap annotated	§8.3
<code>fig3_holdout_tawos.svg</code>	Held-out P/R/F1 across three approaches	§8.4
<code>fig4_ablation.svg</code>	Per-risk-type F1, graph vs naive	§8.6

Design decisions worth mentioning if asked: the palette was *validated*, not eyeballed — a blue/red/green candidate failed CVD checking at protan ΔE 3.4, and the shipped palette passes at ΔE 23.7. Because IEEE papers are printed B&W, and blue vs purple collapse in grayscale (ΔL 0.018), Fig 3 carries hatch patterns and direct value labels as secondary encoding.

Ready-to-paste IEEEtran captions and `svg→pdf` conversion commands are in `ai-service/eval/figures/README.md`. Use `figure*` (starred) environments for full width.

11. Related work positioning

Position against three groups:

Group	Examples	What they lack
Traditional PM tools	Jira, Linear, Asana, Monday.com	Surface current state, not predictive evidence-linked risk; "at risk" flags are opaque
LLM multi-agent systems	AutoGPT, MetaGPT, ChatDev, AgentBench	Optimize task completion by LLM agents; no deterministic, checkable risk computation
XAI in software engineering	SHAP/LIME defect prediction, explainable effort estimation	Explain a learned model's prediction. Here the computation is a deterministic graph invariant needing no post-hoc explanation; the LLM's role is narration under a citation constraint

The differentiating combination — no comparator has all five: specialist multi-agent decomposition + graph-computed (not LLM-inferred) risk + mechanically enforced evidence grounding + deterministic fallback with identical schema + held-out real-world evaluation including reported negative results.

12. IEEE paper structure mapping

Suggested 8–12 pages for IEEE Access; 6–8 for a conference.

§	Section	Source in this report	Pages
—	Abstract	§1 + the one-sentence novelty statement in §4	200 words
I	Introduction	§2 problem → §11 gap → §3 approach → §4 contributions	1–1.5
II	Related Work	§11 (build the 3-group table)	1–1.5
III	System Architecture	§5 + Fig. 1	1.5–2
IV	Graph-Grounded Prompting ← core	§4 Novelty 1–3, §5.3–5.5, §8.10 bugs as rigor	2–3
V	Multi-Agent Design	§5.6 + AgentOutput schema	1–1.5
VI	Experimental Setup	§7 datasets, split discipline, point-in-time reconstruction	1
VII	Results	§8.2–8.7 + Figs. 2, 3, 4	2–3
VIII	Discussion	§8.5 negative results, §4 Novelty 4, ethics of risk-scoring people	1

§	Section	Source in this report	Pages
IX	Threats to Validity	§9 — reproduce honestly, all ten	0.5
X	Conclusion & Future Work	§1 + §14	0.5

Tone guidance. The strongest defensible claim is:

"A simple, explainable, zero-parameter graph rule matches or beats a tuned learned model on unseen projects, at 487× lower prompt cost."

Do not overclaim beyond that. Report the negative results (§8.5) **prominently** rather than burying them — they are what makes the headline credible.

Ethics paragraph worth writing: risk-scoring *people* (workload, silent-member detection) carries surveillance and fairness concerns. A paragraph acknowledging this strengthens the paper.

13. Reproducibility

```
cp .env.example .env          # GROQ_API_KEY optional – full fallback mode works without it
make up && make db-migrate && make db-seed
make test                     # backend 29/29, ai 43/43, frontend 27/27

python -m eval.run_eval        # synthetic detection, tokens, latency, grounding
python -m eval.boundary_eval   # threshold-boundary sensitivity
python -m eval.datasets.tawos_load # TAWOS .sql -> SQLite (~4.3GB, one pass)
python -m eval.datasets.tawos_split # fixed project-level split
python -m eval.datasets.tawos_score # all-987-sprint result
python -m eval.datasets.tawos_holdout_eval # the one-shot held-out number
python -m eval.ablation_naive_vs_graph # naive-vs-graph ablation [needs API key, billed]
python -m eval.live_llm_eval      # schema validity + Cohen's κ [needs API key, billed]
python -m eval.make_figures       # regenerate all 4 figures (offline)
```

Required citation: Tawosi, V., Al-Subaihini, A., Moussa, R., & Sarro, F. *A Versatile Dataset of Agile Open Source Software Projects*. MSR 2022. doi:10.1145/3524842.3528029

14. What is still missing

Honest gap list, in priority order.

Prio rity	Gap	Effor t	Why it matters
1	Bootstrap confidence intervals on held-out F1	~half day	Turns one point estimate into a defensible interval — the most likely reviewer request
2	Bibliography (~15–20 refs)	1 day	Currently comparators are named but not cited
3	IEEEtran conversion + svg→pdf	1 day	Submission format
4	Human evaluation of explanation quality	1–2 week s	The biggest unmeasured claim for an explainability paper
5	Grounding ablation (run with the check disabled)	~1 day	Proves the safety net catches something real, not just synthetic fabrications
6	Balanced live-LLM sample (paid tier)	days	Fixes n=1 categories in the κ breakdown
7	Validate remaining 5 risk types on real data (AMI/Enron)	mont hs	Removes the largest scope limitation

Verdict on readiness:

Venue	Ready?
Regional IEEE conferences (ICCCNT, ICACCS, etc.)	Yes, comfortably above bar
IEEE Access	Yes, after items 1–3
IEEE ICSME / SANER	Borderline — add items 4–5
IEEE TSE / ICSE / ASE	No — needs items 4–7

License: MIT. Third-party: Groq, NetworkX, FastAPI, SQLAlchemy, Pydantic, Next.js, shadcn/ui, and the TAWOS dataset (Apache 2.0, MSR 2022 — citation required).